

CSE 256 PA2 Report - Transformers

Atharv Sunil Biradar
Electrical and Computer Engineering
University of California San Diego
abiradar@ucsd.edu

1 Part 1 : Encoder Only Classifier Model

1.1 Sanity Check for Encoder

In this part of the project, the encoder architecture was implemented, and attention maps were extracted from the final attention block to analyze how each token in the input sequence attends to others during the encoding process. As illustrated in Figure 1, these attention maps provide valuable insights into the intricate relationships between tokens. The observed attention patterns reveal that certain tokens engage in strong interactions, while others exhibit weaker connections. Notably, only a select few tokens show significant correlations with their neighbors; in the attention map, darker shades represent weaker connections, whereas lighter shades highlight stronger associations.




Figure 1: Attention Maps for Encoder

This selective attention mechanism demonstrates the transformer model's ability to focus computational resources on specific tokens that are essential for capturing context and semantics. By prioritizing these pivotal tokens, the model effectively encodes and captures complex dependencies and interactions within the text, enhancing its ca-

capacity to interpret and represent textual information with depth and accuracy. This analysis of attention patterns is a central element of the project, offering insight into the model's language processing and informing further architectural and performance improvements.

1.2 Evaluation

Table 1 presents the training and evaluation results for the encoder-only classifier model across 15 epochs. The table tracks the model's training loss, training accuracy, and test accuracy, highlighting the progression of learning and performance over time.

In the initial epochs, the model demonstrates rapid improvements, with training loss decreasing from 1.392 at epoch 1 to 0.790 by epoch 4, accompanied by a substantial rise in training accuracy from 45.41% to 83.94%. Correspondingly, test accuracy increases from 34.00% to 70.27%, indicating that the model is beginning to generalize well to unseen data.

As training continues, the model achieves further accuracy gains, reaching a peak training accuracy of 99.86% by epoch 15. However, test accuracy begins to plateau around 83% from epoch 10 onward, peaking at 84.00% in epoch 10 before stabilizing. This trend suggests that while the model continues to perform exceptionally well on the training data, improvements in generalization have slowed, potentially indicating a slight over-fitting tendency as training accuracy approaches saturation.

Overall, the encoder-only classifier model demonstrates effective learning, with notable improvements in both training and test accuracies, especially in the initial epochs. The stabilization of test accuracy after epoch 10, however, suggests a diminishing return on additional epochs for test set performance. This evaluation supports the model's strength in encoding and classifying within the given dataset, although further tuning or regularization could potentially enhance generalization in future iterations.

Epoch	Train Loss	Train Acc (%)	Test Acc (%)
1	1.392	45.411	34.00
2	1.129	58.891	50.667
3	0.966	75.478	60.933
4	0.790	83.939	70.267
5	0.567	89.149	74.80
6	0.306	94.312	77.867
7	0.551	94.503	80.133
8	0.325	98.327	82.80
9	0.171	98.996	83.20
10	0.242	99.474	84.00
11	0.109	99.570	83.20
12	0.204	99.761	83.067
13	0.017	99.761	83.867
14	0.032	99.761	83.467
15	0.023	99.857	83.467

Table 1: Results for Encoder Training

2 Part 2 : Decoder Only Language Modeling

Similar to the encoder, I extracted the attention maps from the last attention block of the decoder, as shown in Figure 2. The upper triangular matrix of each map is dark due to masked attention, which restricts each token to attend only to preceding tokens.




Figure 2: Attention Maps for Decoder

2.1 Evaluation

In Tables 2 and 6, the perplexity values of the decoder are presented, providing insights into model performance across training iterations and on specific test datasets.

Table 2 reports the perplexity values observed every 100 iterations during training, starting from an initial perplexity of 6672.14 at iteration 0 and demonstrating a marked decrease over time. By

iteration 500, the perplexity value reaches 195.70, indicating that the decoder is learning effectively and achieving better predictive accuracy as training progresses. The steep reduction from initial to final perplexity reflects the model’s ability to capture patterns within the training data, with diminishing gains in perplexity reduction as training stabilizes around iteration 500.

Iteration	Perplexity
0	6672.14
100	550.47
200	415.45
300	313.02
400	240.07
500	195.09

Table 2: Perplexity Values of Decoder every 100 Iterations

Table 6 provides perplexity values at the 500th iteration for the training dataset and three distinct test datasets (H. Bush, Obama, and W. Bush). While the training set perplexity at this step is 195.70, the test datasets exhibit higher perplexity values, ranging from 404.72 (Obama) to 515.45 (W. Bush). These elevated perplexity values on the test datasets indicate some difficulty in generalizing beyond the training data, which could reflect either a domain shift between training and test datasets or the need for further regularization to enhance generalization.

Dataset - Step 500	Perplexity
Train	195.09
H. Bush	437.80
Obama	404.72
W. Bush	515.45

Table 3: Perplexity Values of Decoder every 100 Iterations

Overall, the decoder exhibits strong learning within the training set, as evidenced by significant perplexity reductions. However, the disparity between training and test perplexities suggests an opportunity for model refinement to improve performance on unseen data.

The observed variation in perplexity scores across datasets can likely be attributed to the following factors:

- **Dataset Characteristics:** Each evaluation dataset, representing distinct figures such as Obama, H. Bush, and W. Bush, contains unique linguistic attributes, vocabulary ranges, and stylistic elements. These distinct

features impact the model’s performance, as it must adapt to unfamiliar language patterns specific to each dataset. Consequently, differences in linguistic complexity or terminology usage across datasets contribute to variability in perplexity. For instance, datasets with more complex sentence structures or rare vocabulary may yield higher perplexity values, indicating increased model difficulty in accurately predicting these tokens.

- **Domain Specificity:** The speeches of different political figures are shaped by factors such as political ideology, target audience, historical context, and even individual speaking styles, resulting in divergent language characteristics. These variations can affect how well the model generalizes, as it may capture language patterns associated with specific figures or periods more effectively than others. If the training data is limited in covering diverse language styles, this may further exacerbate differences in performance across datasets. Consequently, datasets reflecting certain figures’ linguistic nuances or rhetorical styles may pose greater challenges for the model, resulting in elevated perplexity scores.
- **Model Adaptability:** The extent to which the model can generalize to new linguistic contexts depends on its adaptability to unseen data. High perplexity values on datasets outside the training distribution may indicate that the model struggles with language patterns or idioms not encountered during training. This effect may be more pronounced in datasets with highly distinctive or domain-specific language, signaling a potential need for further model refinement, such as domain adaptation or regularization, to improve generalization capabilities.
- **Training Data Distribution:** The composition and distribution of the training data also play a critical role in model performance across evaluation datasets. If the training data heavily represents certain linguistic styles or topics, the model may learn biases that favor these patterns, thereby impacting its effectiveness on datasets with contrasting language styles. Ensuring a diverse training dataset could help mitigate this effect by providing the model with exposure to a broader range of linguistic variations, potentially reducing the observed perplexity disparities.

These factors underscore the importance of considering dataset diversity and domain adaptation when evaluating language models, as variability in

linguistic features and domain specificity can significantly influence model performance across different contexts.

3 Part 3: Architecture Exploration: ALiBi and RoPE

To optimize the Transformer architecture for handling longer sequences and to improve computational efficiency, this project integrates two positional encoding methods: ALiBi (Attention with Linear Biases) and Rotary Positional Encoding (RoPE). This section details the combined approach, the underlying algorithmic principles, and its performance in comparison to traditional positional encoding.

The integrated ALiBi-RoPE encoding leverages the strengths of both techniques, with ALiBi providing a relative positional bias that enables efficient handling of long sequences by modifying attention scores based on token distance. This removes the need for explicit positional embeddings, which reduces memory requirements as sequence length increases, making it particularly advantageous for long-sequence tasks.

RoPE complements this by embedding positional information directly within the attention mechanism through rotational transformations, allowing the model to capture relative positions and dependencies between distant tokens with improved accuracy. By encoding positions within the query and key projections, RoPE enhances the model’s generalization to extended sequences, maintaining robust performance across varying context lengths.

This combined approach allows the model to process lengthy sequences while preserving computational efficiency and minimizing memory overhead. The subsequent sections evaluate the performance of the ALiBi-RoPE encoding, highlighting its effectiveness in retaining long-term dependencies and its suitability for applications requiring extensive context modeling.

3.1 ALiBi Encoding (Attention with Linear Biases)

ALiBi (Attention with Linear Biases) was proposed by Press et al. as a novel positional encoding method that enables transformers to attend effectively to long-range dependencies without the traditional, position-specific embeddings. The core idea is to add a linearly increasing bias to the attention scores based on the distance between tokens, allowing the transformer to handle varying sequence lengths without explicit positional embeddings.

The formulation of ALiBi is as follows:

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} + M \right) V$$

where Q and K are the query and key matrices, d_k is the dimension of the key vectors, and M is a matrix of linear biases based on token distances. Specifically, the bias increases linearly with distance, so more distant tokens contribute less to the attention score. This bias helps reduce computational overhead and improves attention on relevant tokens without explicit encoding, especially useful in tasks with variable-length sequences.

3.2 Rotary Positional Encoding (RoPE)

Rotary Positional Encoding (RoPE) is a method designed to incorporate positional information in Transformer models to better handle long sequences. Unlike traditional absolute positional encodings, which may struggle with generalization to unseen sequence lengths, RoPE encodes relative positional relationships directly within the model’s attention mechanism. This is achieved by applying a rotational transformation to the query and key vectors based on each token’s position, thereby capturing relative distances between tokens. As a result, RoPE facilitates robust handling of longer contexts, making it well-suited for tasks that require extensive context modeling.

3.2.1 Algorithmic Foundations

The algorithmic foundation of RoPE involves modifying the query and key vectors in the attention mechanism to embed relative position information. For each token position, RoPE applies a rotation matrix to the query and key vectors. This rotational transformation ensures that the similarity scores between tokens depend on their relative positions within the sequence, rather than absolute positions. Specifically, by encoding each position as a rotation within the attention inner product, RoPE effectively represents positional relationships without the need for explicit position embeddings. This approach enhances the model’s flexibility in managing variable sequence lengths and capturing dependencies across tokens with high computational efficiency.

Mathematically, given a position p and a query or key vector x , RoPE applies a rotation \mathbf{R}_p to x as follows:

$$\mathbf{R}_p(x) = x \cdot \cos(p) + \text{rotate}(x) \cdot \sin(p)$$

where $\text{rotate}(x)$ applies a transformation to x that rotates its components by 90 degrees (for example, by rearranging the dimensions of x). This rotation allows the inner product between the query and key vectors to encode relative position information, enabling efficient and adaptive handling of long-range dependencies in sequence data.

3.3 Implementation

3.3.1 Implementation Details

In this implementation, the model’s attention mechanism is modified to incorporate both AliBi and RoPE. AliBi introduces a linear bias to the attention scores based on token distance, allowing the model to prioritize nearby tokens and gradually reduce attention on distant tokens. This bias is applied directly to the attention weights, avoiding the need for explicit positional embeddings and thereby reducing memory usage, especially for long sequences.

RoPE, on the other hand, introduces rotational transformations to the query and key vectors, encoding relative positional relationships within the inner product of these vectors. By rotating the query and key vectors according to their positional offsets, RoPE enables the model to capture dependencies between tokens based on their relative positions. This is beneficial for representing long-range dependencies while keeping memory requirements minimal. The combined approach leverages the strengths of both AliBi’s linear biases and RoPE’s rotational encodings, enhancing the model’s capacity for context retention and adaptation to various sequence lengths.

3.3.2 Advantages in Terms of Execution

The integration of AliBi and RoPE brings notable advantages in terms of execution efficiency:

- **Memory Efficiency:** By removing the need for explicit positional embeddings, AliBi significantly reduces memory consumption, particularly for long sequences. This allows the model to handle extended sequences without a substantial increase in memory usage.
- **Improved Long-Range Attention:** RoPE’s rotational transformations enable the model to capture long-range dependencies more effectively by encoding relative positions within the query and key vectors. This improves the model’s ability to understand the structure and relationships in lengthy sequences.
- **Reduced Computational Complexity:** The combination of AliBi and RoPE allows for efficient scaling with sequence length. AliBi’s linear biases reduce the emphasis on distant tokens, while RoPE maintains attention stability across varying contexts, leading to lower computational complexity and faster execution times.
- **Enhanced Flexibility for Variable-Length Sequences:** The AliBi-RoPE implementation supports variable-length inputs without requiring re-training for

different sequence lengths. This adaptability is crucial for tasks with dynamic sequence lengths, such as document summarization or real-time processing of streaming data.

Overall, the combined use of AliBi and RoPE enables the Transformer to maintain high performance on long-sequence tasks while optimizing memory and computational requirements, making it an effective solution for efficient context modeling.

3.4 Observation and Analysis: With AliBi and RoPE

The integration of AliBi and RoPE into the Transformer architecture resulted in noticeable improvements in both the classification and language modeling tasks. This section highlights the observations derived from the training and test metrics and presents a comparative analysis of performance with and without the use of these positional encoding techniques.

Epoch	Train Loss	Train Accuracy (%)	Test Accuracy (%)
0	0.877	51.291	40.667
1	0.897	65.966	54.267
2	0.719	72.992	63.680
3	1.139	87.094	72.933
4	0.638	92.878	78.267
5	0.368	95.793	82.933
6	0.451	97.228	82.933
7	0.257	98.279	82.933
8	0.073	99.570	84.000
9	0.022	99.904	84.000
10	0.013	99.904	84.667
11	0.010	99.904	84.667
12	0.016	99.904	84.333
13	0.036	99.522	83.867
14	0.003	99.904	84.133

Table 4: Training and Testing Metrics across Epochs

3.4.1 Training and Testing Metrics (Classification Task)

The training and testing metrics, as shown in Table 4, exhibit improved performance with the use of AliBi and RoPE. The following observations can be made:

- **Training Loss:** The training loss starts at 0.877 at epoch 0 and decreases steadily, reaching a minimal 0.003 by epoch 14. The presence of AliBi and RoPE allowed the model to converge more efficiently during training, as the

positional encodings provided better context preservation for long-range dependencies.

- **Training Accuracy:** The model achieved a very high training accuracy, peaking at 99.904% by epoch 10. The use of positional encodings significantly improved the model’s ability to capture intricate patterns in the data, allowing it to effectively learn from the training set.
- **Test Accuracy:** The test accuracy increased more rapidly compared to the model without AliBi and RoPE, reaching 84.000% at epoch 10. The introduction of these techniques reduced overfitting and helped the model generalize better, thus enhancing performance on unseen data.

3.4.2 Perplexity (Language Modeling Task)

Table 5 shows the training perplexity values for the model with AliBi and RoPE. The following points highlight the impact of these techniques on perplexity during training:

- **Perplexity Decrease:** From the initial value of 6474.0522 at iteration 0, the perplexity steadily decreases to 195.8733 at iteration 500. The use of AliBi and RoPE accelerates the learning process by enabling the model to better capture long-term dependencies, resulting in a more rapid reduction in perplexity compared to the model without these techniques.
- **Fluctuations in Perplexity After Iteration 500:** The perplexity increases to 454.7881 at iteration 600, similar to the non-AliBi and RoPE model, indicating that the model may be starting to overfit. This suggests that further fine-tuning or regularization may be needed for optimal performance.

Iteration	Train Perplexity
0	6474.0522
100	586.1444
200	467.2596
300	343.6937
400	251.1837
500	195.8733
600	454.7881

Table 5: Training Perplexity across Iterations

Table 6 presents the perplexity values for different datasets after 500 iterations. The following observations can be made:

- **Dataset-Specific Performance:** The model with AliBi and RoPE shows lower

perplexity on the training dataset (195.873) compared to datasets such as H. Bush (454.79), Obama (405.92), and W. Bush (512.52). While the model performs well on the training set, the perplexity for certain test datasets remains relatively higher, reflecting the challenges of generalizing across varied linguistic styles.

Dataset - Step 500	Perplexity
Train	195.873
H. Bush	454.79
Obama	405.92
W. Bush	512.52

Table 6: Perplexity Values of Decoder every 100 Iterations

4 Comparative Analysis: With and Without AliBi and RoPE

The introduction of AliBi and RoPE positional encoding techniques significantly altered the performance of the Transformer model in both classification and language modeling tasks. Below is a comparative analysis based on the metrics from the previous section.

4.1 Classification Task

When compared to the model without AliBi and RoPE, the Transformer model with these encoding techniques displayed:

- Faster Convergence:** With AliBi and RoPE, the model converged faster, achieving higher training and test accuracy more rapidly. The model’s ability to generalize on unseen data (test accuracy) improved, as evidenced by the test accuracy reaching 84.000% at epoch 10, compared to lower performance without these techniques.
- Better Generalization:** The introduction of AliBi and RoPE helped the model generalize better across different datasets, as seen from the test accuracy. Without these techniques, the model experienced overfitting, with a larger gap between training and testing performance.

4.2 Language Modeling Task

The language modeling task also benefited from the use of AliBi and RoPE, as shown in Table 5 and Table 6:

- Reduced Perplexity:** The perplexity values in the training phase were considerably lower

with AliBi and RoPE, indicating better performance in predicting the next token. Without these positional encodings, the perplexity remained high, suggesting difficulty in capturing long-range dependencies.

- Improved Handling of Long Sequences:** The models with AliBi and RoPE were able to more effectively handle long-range dependencies, which was particularly important for processing the sequential nature of language data. Without these techniques, the model’s ability to capture distant relationships in the data was hindered, leading to higher perplexity values.
- Improved Generalization Across Datasets:** The model with AliBi and RoPE showed better stability across datasets, with lower perplexity in the training set and slightly higher but still improved perplexity on test datasets like H. Bush, Obama, and W. Bush. This indicates that the model is better equipped to generalize and maintain performance on unseen data.

5 Conclusion

This paper explores the use of advanced positional encoding techniques, AliBi and RoPE, to enhance the Transformer architecture for classification and language modeling tasks. Both methods demonstrated improved model performance by capturing long-range dependencies more effectively, reducing perplexity, and improving generalization.

The integration of AliBi and RoPE resulted in faster convergence and better test accuracy compared to the baseline model. These techniques also proved more scalable for handling longer sequences, offering significant computational benefits.

In conclusion, AliBi and RoPE enhance Transformer models’ efficiency and accuracy, making them valuable for a range of natural language processing tasks. Future research could explore further optimizations of these techniques and their applicability to other deep learning models.